

High-Level Design (HLD): Healthcare PDF Hub (Modular)

1) Purpose & Goals

Purpose. A Streamlit application that lets users upload, browse, and retrieve knowledge from healthcare PDFs (Medical Documents, Medicine Details, Hospital Details). It adds lightweight RAG (retrieve-and-generate) using FAISS + sentence-transformer embeddings and an LLM (Euri AI wrapper) to answer questions grounded in the uploaded content.

Primary goals

- Organize PDFs by domain with quick preview/download.
- Read PDFs from local “resource folders” (Windows paths with relative fallbacks).
- Build in-memory search over uploaded PDFs (chunk → embed → FAISS).
- Ask questions and get answers that cite the retrieved text.
- Ship as a modular codebase that’s easy to extend.

Non-goals (current version)

- No multi-user tenancy or persistent database.
- No server-side file persistence (uploads live in session only).
- No OCR for scanned PDFs (image-only pages won’t extract text).
- No enterprise auth/SSO.

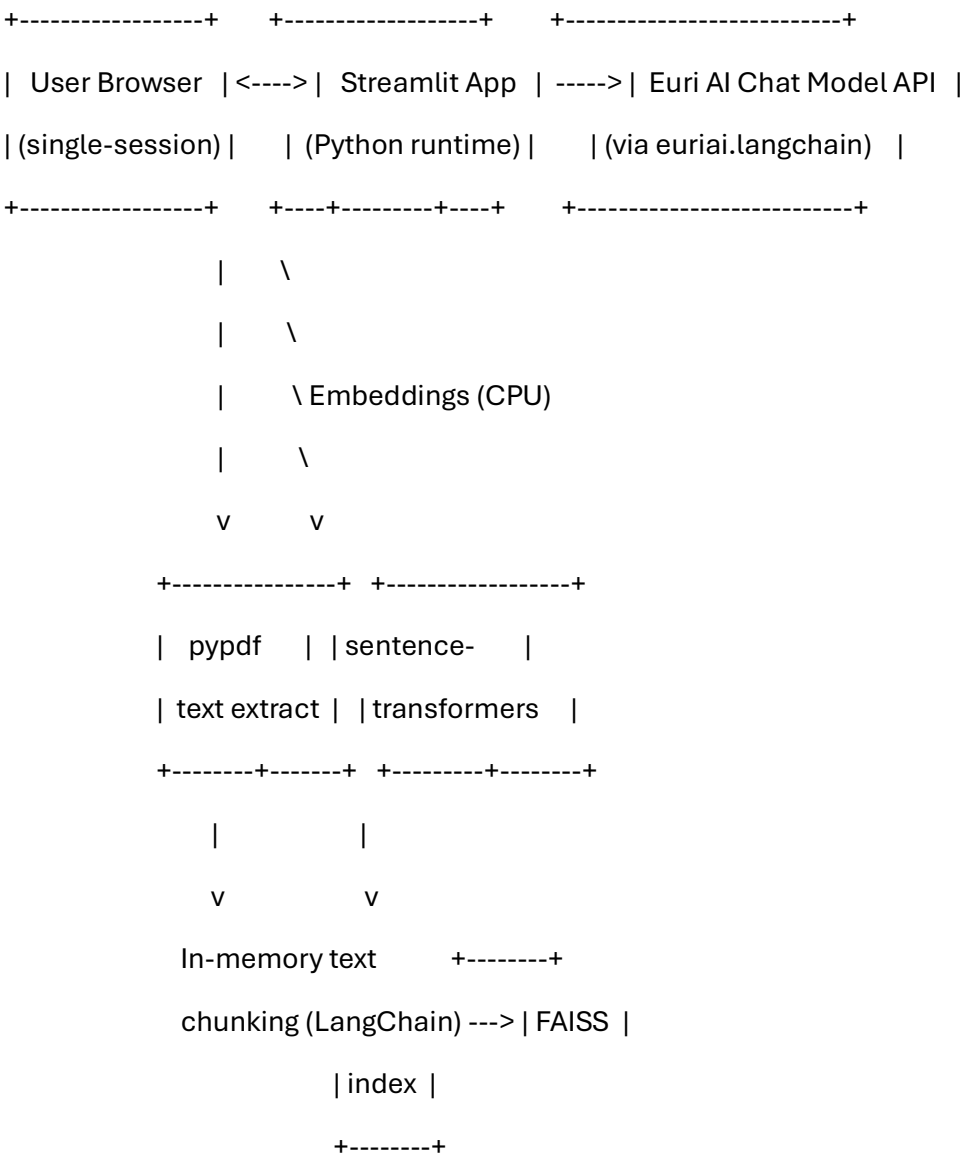
2) Users & Use Cases

- **Clinician/Pharmacist:** Search guidance leaflets, dosage notes, hospital brochures.
- **Patient/Patient Advocate:** Organize personal medical reports; ask clarifying questions.
- **Ops/Admin:** Maintain curated “resource folders” and export bundles (ZIP).

Key flows:

1. Upload PDFs → Add to Library → Preview/Download.
 2. Build FAISS index from uploaded files → Query (prompt) → Answer with citations.
 3. Browse local resource folders and one-click **Download ALL (ZIP)** per folder.
-

3) System Context & Architecture



Key external libs

- streamlit (UI), pypdf (text extraction), sentence-transformers + faiss-cpu (vector store), langchain(-community, -text-splitters) (splitters & vectorstore wrapper), euriai.langchain (LLM wrapper), python-dotenv (env).

4) Module View

Top-level layout

Healthcare-PDF-Hub-Modular/

```

└─ app.py
└─ requirements.txt
└─ README.md
└─ src/
    └─ healthcare_pdf_hub/
        └─ __init__.py
        └─ catalogs.py    # medicine catalog/brands, hospitals list
        └─ config.py      # resolves resource folders (env → abs → rel)
        └─ ui/
            └─ components.py # reusable Streamlit UI pieces
        └─ utils/
            └─ pdf_utils.py # pypdf helpers, preview HTML, folder listing, ZIP
            └─ chat_model.py # Euri AI wrapper (create & invoke)
            └─ faiss_utils.py # create_faiss_index, retrieve_relevant_docs

```

Key components

- **app.py**
 - Sets page, tabs, and orchestrates flows for each tab.
 - Initializes chat model via .env (EURI_API_KEY) with @st.cache_resource.
 - Uses session_state (uploads, per-tab last batch, vectorstores) to manage in-memory state.
- **utils/pdf_utils.py**
 - extract_text_from_pdf(bytes), get_page_count, pdf_preview_html, list_pdfs_from_folder(Path), make_zip_from_items.
- **utils/faiss_utils.py**
 - create_faiss_index(texts) → FAISS.from_texts with HuggingFaceEmbeddings (e.g., all-MiniLM-L6-v2).
 - retrieve_relevant_docs(vectorstore, query, k=4).
- **utils/chat_model.py**

- `get_chat_model(api_key) → euri.ai.langchain.create_chat_model`.
 - `ask_chat_model(chat_model, question) → .invoke()`.
 - **catalogs.py**
 - `MEDICINE_CATALOG, MEDICINE_BRANDS, HOSPITALS_2025`.
 - **config.py**
 - `choose_resource_dirs()` uses env vars `HPDFHUB_MEDICAL_DIR`, `HPDFHUB_MEDICINE_DIR`, `HPDFHUB_HOSPITAL_DIR`; otherwise absolute Windows defaults; then relative fallbacks.
 - **ui/components.py**
 - `process_uploads(files, bucket_key)` stores bytes in `st.session_state.uploads`.
 - `render_bucket_table(bucket)` shows a small table + preview expander.
-

5) Data Design

Session state keys

```
st.session_state = {
    "uploads": {
        "medical": [ {name, size, pages, uploaded_at, data}, ... ],
        "medicine": [ ... ],
        "hospital": [ ... ]
    },
    "medical_last_batch": [ {name, data}, ... ],
    "medicine_last_batch": [ {name, data}, ... ],
    "hospital_last_batch": [ {name, data}, ... ],

    "medical_vectorstore": <FAISS>,
    "medicine_vectorstore": <FAISS>,
    "hospital_vectorstore": <FAISS>,
}
```

```
# optional brief notes if retained in some tabs

"hosp_notes": { "<Hospital — City>": "<text>" },

"med_notes": { "<MedicineName>": "<text>" },

}
```

In-memory vector index

- Texts are **chunked** with RecursiveCharacterTextSplitter (default: 1,000 chars, 200 overlap).
 - Embeddings via HuggingFaceEmbeddings, model sentence-transformers/all-MiniLM-L6-v2 (configurable).
 - FAISS index held in memory; not persisted by default.
-

6) Core Flows (Sequence)

A) Upload & Library

1. User selects PDFs with st.file_uploader.
2. Click **Add to Library** → process_uploads() stores bytes in session_state.uploads[tab].
3. Also cache the **last batch** for that tab (e.g., medicine_last_batch).

B) Build Index & Ask

1. User enters **prompt** (disabled until Library has docs) or selects entity (medicine/hospital).
2. Extraction: extract_text_from_pdf for each file in last batch (or all library items).
3. Chunking: splitter.split_text.
4. Embedding + Index: create_faiss_index(chunks) → store in session_state[tab_vectorstore].
5. Retrieval: retrieve_relevant_docs(vectorstore, prompt_or_selection).
6. Compose **system prompt** with context; call ask_chat_model.
7. Render answer; optionally show retrieved snippets.

C) Resource Folders

1. Read PDFs from configured directories via list_pdfs_from_folder.

2. Offer **per-file downloads** and a **Download ALL (ZIP)** button per folder (make_zip_from_items).
-

7) Functional Requirements

- **FR1:** Upload PDFs per domain and show them in a Library table.
 - **FR2:** Preview PDFs inline and download individually.
 - **FR3:** Read PDFs from local resource directories, list, and download (including ZIP-all).
 - **FR4:** Build a searchable index from uploaded PDFs and run queries.
 - **FR5:** Answer user questions with LLM using retrieved context; keep answers conservative and cite sources (as available in retrieved chunks).
 - **FR6:** Disable prompts until docs are added to Library.
-

8) Non-Functional Requirements

- **Performance:**
 - Reasonable on CPU (MiniLM embeddings).
 - Chunking runs client-triggered; document size dependent.
- **Scalability:**
 - Single session; in-memory state; suitable for desktop/single user or small teams behind a shared session.
- **Reliability:**
 - Handle empty/invalid PDFs; show warnings.
- **Security/Privacy:**
 - API key from .env.
 - No server-side persistence of uploads by default; sensitive medical data stays in session memory unless user changes code to persist.
 - **Do not log extracted text** to console in production.
- **Compliance:**
 - App is informational; not a medical device. Answers stress safety disclaimers.

- **Observability:**
 - Streamlit logs and minimal UI alerts; optional future: structured logging.
-

9) Configuration

- **Environment variables**
 - EURI_API_KEY (required) — for chat model.
 - HPDFHUB_MEDICAL_DIR, HPDFHUB_MEDICINE_DIR, HPDFHUB_HOSPITAL_DIR (optional) — override resource folders.
 - **Requirements**
 - streamlit, pypdf, python-dotenv,
 - langchain, langchain-community, langchain-text-splitters,
 - sentence-transformers, faiss-cpu, torch (CPU or CUDA),
 - euriai (or the package that provides euriai.langchain).
-

10) Error Handling & Edge Cases

- **Scanned PDFs:** pypdf may return empty text → show warning; optional roadmap: OCR (e.g., pdf2image + pytesseract).
 - **Model init failure:** show visible error if EURI_API_KEY missing or SDK unavailable.
 - **Large files:** Streamlit upload limit (default 200MB/file). Chunk size tunable to balance recall vs. speed.
 - **Package conflicts:** Pin compatible versions (Torch ↔ Transformers ↔ Sentence-Transformers ↔ Accelerate).
-

11) Security Considerations

- Do not store or transmit PHI beyond the user's control; avoid server persistence unless secured.
- Keep .env out of version control.
- Sanitize/limit what is sent to the LLM (only retrieved context, not entire documents).

- Consider content filtering and PII redaction for production.
-

12) Extensibility Roadmap

- **Persistence:** Save FAISS indexes & metadata to disk (e.g., `vectorstore.save_local()`).
 - **OCR:** Add OCR fallback path with a toggle per tab.
 - **Source citations:** Attach metadata={"source": file_name, "page": n} to chunks; surface in answers.
 - **Better RAG:** Use MultiQueryRetriever / Hybrid (BM25 + dense) retrieval.
 - **Multi-user:** User auth; isolate session stores; server-side storage (S3/GCS/Azure).
 - **Analytics:** Query history, feedback thumbs (store locally or in a DB).
 - **Prompting:** Centralize prompt templates with guardrails.
-

13) Deployment

- **Local/dev:**
 - `pip install -r requirements.txt`
 - `streamlit run app.py`
 - **Server:**
 - Reverse proxy (Nginx) → Streamlit.
 - Secure environment variables.
 - Optional: mount volumes for resource folders and persistent indexes.
-

14) Testing Strategy

- **Unit:** `pdf_utils.extract_text_from_pdf`, `make_zip_from_items`, `faiss_utils.create_faiss_index`.
- **Integration:** End-to-end on small sample PDFs per tab; verify retrieval returns expected chunks.
- **UX:** Prompt disabled until Library has docs; "Download ALL" includes all files.

15) Open Risks / Assumptions

- Embedding/model downloads on first run (network/time). Consider packaging or pre-warming.
- Session memory growth with many/large PDFs. Option: cap, or paginate Library.
- Legal/clinical risk: reinforce disclaimers; never present guidance as medical advice.

16) Quick Data Flow (per tab)

1. **Upload → Add to Library** → store {name, data, meta} in `st.session_state.uploads[tab]` and `*_last_batch`.
2. **Submit/Run Search** → extract text (pypdf) → chunk (LangChain) → embed (MiniLM) → index (FAISS) → retrieve (K docs).
3. Compose **LLM prompt** with retrieved context + user question → `ask_chat_model` → UI renders **Answer** + (optional) retrieved snippets.